

Reveal AI: Chat Therapist and Intent Detector

Version: 0.3

Author: Ahmed Sheikh

Date: 5 August 2026

Delivery date: 26 August 2026

Status: Draft for client review

Read this part first

Three weeks is the deadline, so this document is built backwards from that date. Some things you asked for are in. Some are moved to a second phase. Section 8 lists exactly what got cut and why.

I have made assumptions where I did not have an answer. They are in Section 9. If any are wrong, tell me and I will adjust, but some of them are load-bearing for the date.

Section 7 has five decisions I need from you. I need them within 48 hours of you reading this. Every day past that moves the delivery date by a day, because

the work in weeks two and three depends on those answers.

There is a companion demo plan. Six short calls across the three weeks, each showing you a working piece rather than a status update.

1. What we are building

Reveal AI already has a 60-second daily voice check-in. This project adds two things on top of it.

Chat Therapist. A conversation feature that pulls from the user's past check-ins, so it knows what they have been dealing with instead of starting cold every time.

Intent Detector. The user records their own voice once. Then they record a conversation, pick a scenario, and get an analysis of how the other person came across.

There is also a bug to clear first. Chat and profile data is not saving right now. That gets fixed in the first two days because nothing else can be built on top of it.

2. In scope

- Persistence fix

- Chat therapist that reads mood and session history
- Profile History screen with a mood chart
- Crisis detection and escalation in the chat
- One-time voice enrolment
- Conversation recording and upload, web browser
- Speaker separation using the enrolled voice sample
- Three scenarios: date, interview, general
- Analysis output stored and viewable later
- Consent screen before recording
- Encryption, user-initiated deletion, and upload failure recovery

3. Out of scope

- Native mobile apps
 - Team tier and manager dashboard
 - Live analysis during a conversation
 - Changes to the existing burnout feature
 - Voice output
 - Languages other than English
 - Legal documents and legal review
-

4. Chat Therapist

Users check in daily but have nowhere to talk. The check-in captures mood and then does nothing with it. Someone who has been sliding for two weeks gets no sign that anyone noticed.

ID	Requirement	How we know it works
T-1	Chat messages save to the database against the user account	Send a message, close the browser, reopen. It is still there.
T-2	Opening a chat loads the last five session summaries and the mood trend	Confirmed by inspecting the payload sent to the model
T-3	Responses reflect that history where it is relevant	A user with three declining sessions gets a reply that acknowledges the pattern
T-4	Ending a session writes a record with summary, mood value, and topics	Row appears within 30 seconds

ID	Requirement	How we know it works
T-5	Profile History lists past sessions newest first with date, mood, summary	Loads under 2 seconds with 100 sessions
T-6	Profile History shows mood over time as a chart	Renders for any user with 2 or more sessions
T-7	User messages are screened for crisis signals before a reply is generated	Defined test phrases trigger escalation, not a normal reply
T-8	Crisis signals interrupt the normal flow and surface support resources	Verified against the same test set

T-7 and T-8 stay in. A product built around emotional wellbeing with no crisis handling is a risk to your users and to you. If you want them removed I need it in writing, and I will say on the record that I advised against it.

5. Intent Detector

ID	Requirement	How we know it works
I-1	User records a voice sample of at least 10 seconds during setup	Sample saves and is retrievable
I-2	User picks date, interview, or general before recording	Selection saves with the session
I-3	A consent screen appears before recording and needs explicit confirmation	Recording will not start without it
I-4	The system splits the recording by speaker and identifies which segments are the enrolled user	Two-person recording in a quiet room, correct on 90 percent of segments or better
I-5	The system analyses the other speaker's tone and word choice against a rubric for the chosen scenario	Output matches a fixed schema and saves against the session

ID	Requirement	How we know it works
I-6	Results are viewable later from history	Still there after logout and login
I-7	Low separation confidence shows a quality warning instead of a result	Tested with a deliberately noisy recording

On accuracy. Speaker separation gets worse as the room gets louder. A restaurant is harder than an interview room, and a date in a restaurant is the hardest common case. I-7 exists so the product says it could not read the recording rather than inventing a confident answer from bad audio.

I would rather ship something that sometimes declines than something that always answers and is sometimes wrong. That is a product position, not just an engineering one, so it is yours to overrule.

6. Data handling and reliability

These are in scope. Two of them are legal obligations, not preferences.

ID	Requirement	How we know it works
N-1	Voice samples and recordings are encrypted at rest and in transit	Verified against stored objects and network traffic
N-2	Users can delete their voice sample and any recording from inside the product	Deleted audio is gone from storage, not just hidden from the interface
N-3	A failed upload or failed processing job does not lose the user's recording	Interrupt an upload mid-transfer, confirm recovery
N-4	A 20-minute recording completes analysis within 3 minutes	Timed against real recordings

N-2 is a right of access and erasure obligation under GDPR and the UAE Personal Data Protection Law. Voice recordings are biometric data. This is not optional and it is not a phase two item.

N-3 is here because the recordings this product captures are often unrepeatable. A user cannot re-record a job interview that already happened.

Automatic expiry of old recordings after the retention period is phase two. Manual deletion covers the obligation in the meantime.

7. Decisions I need from you

D-1. Percentages or observations.

The site currently shows things like Genuine Interest 92 percent. The alternative is plain observations, for example: gave short answers whenever work came up.

I recommend observations. The analysis underneath is a judgement, not a measurement, and running the same recording twice can produce different numbers. A user who notices that will not trust anything else in the product. Percentages look stronger in marketing and match what you have already published. Your call. It changes how I build the output layer, so I need it in week one.

D-2. Consent from the second person.

I have built it in as required. Recording someone without consent carries real legal weight in the UAE. Confirm whether you have had this looked at by a lawyer.

D-3. Who is writing the privacy policy and terms, and by when.

The links on the site currently go nowhere. Voice recordings count as biometric data under both GDPR and the UAE data protection law, and the Intent Detector records people who are not your users and have not agreed to anything. This is not a paperwork issue for me. It decides what the consent screen says and how deletion behaves.

D-4. How long do we keep voice recordings.

Give me a number in days. I have assumed 90.

D-5. Is the Team tier coming.

If manager-facing analysis of employee voice is planned, say so now. It changes the data model, and it carries a very different risk profile from the consumer version.

8. Three-week plan

15 working days. Delivery 26 August.

Week 1, 6 to 12 August. Codebase audit. Persistence fix shipped by day 2. Chat therapist core running against real session history. Decision D-1 locked.

Week 2, 13 to 19 August. Profile History screen and mood chart. Crisis detection and escalation. Voice

enrolment. Speaker separation working end to end on real recordings.

Week 3, 20 to 26 August. Scenario rubrics. Analysis output layer. Confidence thresholds. Encryption, deletion, upload recovery. Edge cases, error states, testing, one revision round.

Chat Therapist is reviewable at the end of week 2 on its own. I would treat that as a checkpoint rather than waiting until the 26th to find out whether it is what you wanted.

Six demo calls are scheduled across the three weeks. Details in the companion demo plan. Each needs about 10 minutes of your time.

What got cut to fit three weeks

- Re-recording your voice sample from settings. Enrol once, support handles resets.
- Automatic deletion of expired recordings. Manual deletion is in, automatic expiry is not.
- Search and filtering on analysis history. List view only.
- More than three scenarios.
- Multi-speaker recordings. Two people only.
- Any accuracy tuning beyond the 90 percent quiet-room threshold.

These are all straightforward to add later. They are out because three weeks does not fit them, not because they are hard.

What puts the date at risk

The persistence bug turning out to be structural rather than a simple fix. I will know by day 2 and will tell you immediately.

Speaker separation missing 90 percent on real recordings. If that happens the options are narrowing the supported scenarios or extending the timeline, and I will bring you the choice rather than quietly shipping something worse.

Late answers on Section 7. Each day of delay is a day on the end.

Scope additions. Anything added after 12 August moves the date.

9. Assumptions

If any of these are wrong the plan changes.

1. Web only.
2. The existing check-in data is usable as the history source for the chat.

- 3. I get full access to the codebase, database, and deployment environment on day 1.
- 4. English only.
- 5. One revision round is included. Further rounds are quoted separately.
- 6. Third-party API costs are billed to you.
- 7. Retention period is 90 days until you tell me otherwise.
- 8. You or your lawyer handle the privacy policy and terms. That work is not in my three weeks and it is not on my critical path, but the feature should not go live without it.

10. Revisions

Version	Date	Change
0.1	4 Aug 2026	First draft, 8 week plan
0.2	5 Aug 2026	Rebuilt for a 3 week deadline. Scope reduced.

Version	Date	Change
0.3	5 Aug 2026	Restored data handling and reliability requirements as Section 6. Encryption, deletion, and upload recovery were dropped in error during the 0.2 compression. Deletion is a legal obligation and should not have been cut.